

Lecture 12

Markov-Chain Monte-Carlo Sampling

Instructor: Shibo Li

shiboli@cs.fsu.edu



Department of Computer Science
Florida State University

- General ideas and Markov chain basics
- Metropolis-Hastings algorithm
- Gibbs sampling
- Hybrid Monte-Carlo

- General ideas and Markov chain basics
- Metropolis-Hastings algorithm
- Gibbs sampling
- Hybrid Monte-Carlo

- Given a probabilistic model

$$p(\mathcal{D}, \mathbf{z}) = p(\mathbf{z})p(\mathcal{D}|\mathbf{z})$$

- How to generate samples from the posterior distribution
(the samples are NOT necessarily independent!)

$$\mathbf{z}_1, \mathbf{z}_2, \dots, \mathbf{z}_N \sim p(\mathbf{z}|\mathcal{D})$$

- Given the posterior samples, what can we do?
- A lot of things
 - Approximate the (marginal) posterior over any subset of variable (unlike message-passing)

$$p(\mathbf{z}|\mathcal{D}) \approx \frac{1}{N} \sum_{n=1}^N \delta(\mathbf{z} - \mathbf{z}_n)$$

- Estimation of any interested statistics/moment

$$\mathbb{E}[f(\mathbf{z})] = \int f(\mathbf{z}) p(\mathbf{z}|\mathcal{D}) d\mathbf{z} \approx \frac{1}{N} \sum_{n=1}^N f(\mathbf{z}_n)$$

- Predictive distribution

$$p(\mathbf{y}^*|\mathcal{D}) = \int p(\mathbf{y}^*|\mathbf{z}) p(\mathbf{z}|\mathcal{D}) d\mathbf{z} \approx \frac{1}{N} \sum_{n=1}^N p(\mathbf{y}^*|\mathbf{z}_n)$$

- Pros
 - Asymptotic convergence to the true posterior (note: deterministic approximation, such as VI, always has discrepancy with the true posterior)
 - Robust to initialization
 - Empirically best and often used as a gold-standard to test other approximate inference algorithms
 - samples are more convenient to use than approximate distributions

- Cons
 - Orders of magnitude slower than VB
 - Hard to diagnosis the convergence
 - Hard for parallelization (sequential sampling nature)
 - Hard for large-scale applications
 - Easily trap into single modes (this is the same as VB)

How to scale up MCMC to big data is a hot research topic!

Sample a sequence of variables using a Markov chain that converges to the desired posterior

$$\mathbf{z}_1 \rightarrow \mathbf{z}_2 \rightarrow \dots \rightarrow \mathbf{z}_n \rightarrow \mathbf{z}_{n+1} \rightarrow \dots$$

$$\mathbf{z}_{n+1} \sim p(\mathbf{z}_{n+1} | \mathbf{z}_n) \quad \lim_{n \rightarrow \infty} p(\mathbf{z}_n) = p(\mathbf{z} | \mathcal{D})$$

Therefore, the MCMC samples are strongly correlated!

- A Markov chain is determined by
 - $p(\mathbf{Z}_1)$: we do not care it much in MCMC sampling
 - Transition kernel: determines the speed of convergence

$$T(\mathbf{z}_n \rightarrow \mathbf{z}_{n+1}) = p(\mathbf{z}_{n+1} | \mathbf{z}_n)$$

if the kernel is the same for all n, the Markov chain is called *homogeneous*

The development of MCMC sampling is the art to design the transition kernel

- What distribution does a MC converge to ?
 - Invariant distribution

$$\int p^*(\mathbf{z}') T(\mathbf{z}' \rightarrow \mathbf{z}) d\mathbf{z}' = p^*(\mathbf{z})$$

We claim that $p^*(\cdot)$ is invariant to the transition kernel T

Also called stationary distribution

Obviously, we want to design a kernel to which the target posterior is invariant

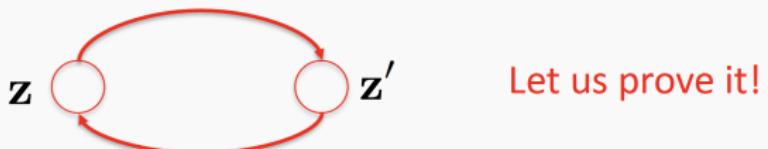
- How to examine invariance?

Sufficient condition (not necessary): *detailed balance*

$$p^*(\mathbf{z})T(\mathbf{z} \rightarrow \mathbf{z}') = p^*(\mathbf{z}')T(\mathbf{z}' \rightarrow \mathbf{z})$$



- How does *detailed balance* lead to invariance?



Let us prove it!

An MC whose stationary distribution and transition kernel respect detailed balance is called *reversible*

- An MC can have multiple stationary distributions; converging to which one depends on $p(z_1)$
- We want our MC only converges to the desired posterior no matter what initial distribution is chosen!
- This property is called ***ergodicity***: *an ergodic MC only converges to one invariant (stationary) distribution*

- Informally, in an ergodic chain, it is possible to go from *every* state to *every* state (not necessarily in one move)
- An ergodic chain is also called *irreducible*
- The invariant (or stationary) distribution of an ergodic chain is called the **equilibrium** distribution

- In MCMC sampling procedure
 - Invariance guarantees the samples will converge to the true posterior (unbiased)
 - Ergodicity guarantees the sample space can be fully explored (rather than partially)
- It can be shown that a homogeneous MC will be ergodic, subject only to weak restrictions on the invariant distribution and transitional kernels

- Conceptually, the sampling contains two stages
 - Before **burn-in**: the MC has yet converged to the invariant distribution. In practice, we usually set up the maximum # of steps before burn-in, and usually various tricks to verify convergence empirically (e.g., look at trace plots).
 - After **burn-in**: the MC has converged. Then we generate the posterior samples. To reduce the strong correlation, we often take every M-th sample (e.g., $M = 5, 10, 20$). We also need to compute the effective sample size (ESS) to ensure the collected samples are enough.

- General ideas and Markov chain basics
- Metropolis-Hastings algorithm
- Gibbs sampling
- Hybrid Monte-Carlo

- A general framework for MCMC

- A general framework for MCMC
- In each step, we first use a proposal distribution to generate a candidate sample, and then decide whether to accept this new sample

- Denote the proposal distribution (not the transition kernel) by $q(\mathbf{z}'|\mathbf{z}_n)$, e.g., $\mathcal{N}(\mathbf{z}'|\mathbf{z}_n, \sigma^2 \mathbf{I})$. Sample the proposal \mathbf{z}' first.
- Accept \mathbf{z}' with probability

$$\min\left(1, \frac{p(\mathbf{z}', \mathcal{D})q(\mathbf{z}_n|\mathbf{z}')}{p(\mathbf{z}_n, \mathcal{D})q(\mathbf{z}'|\mathbf{z}_n)}\right)$$

Unnormalized posterior

Jump back

Jump out

```
graph TD; A[min] -- "Jump back" --> B["1"]; A -- "Jump out" --> C["p(z_n, D)q(z'|z_n)"]; A -- "Unnormalized posterior" --> D["p(z', D)q(z_n|z')"]
```

- Accept \mathbf{z}' with probability

$$\min\left(1, \frac{p(\mathbf{z}', \mathcal{D})q(\mathbf{z}_n|\mathbf{z}')}{p(\mathbf{z}_n, \mathcal{D})q(\mathbf{z}'|\mathbf{z}_n)}\right)$$

Unnormalized posterior

Jump back

Jump out

How do we implement it in practice?

Sample a uniform R.V. u in $[0,1]$, and test if

$$u \leq \exp \left\{ \min \left(0, \log p(\mathbf{z}', \mathcal{D}) + \log q(\mathbf{z}_n|\mathbf{z}') - \log p(\mathbf{z}_n, \mathcal{D}) - \log q(\mathbf{z}'|\mathbf{z}_n) \right) \right\}$$

- If we accept \mathbf{z}'

Set $\mathbf{z}_{n+1} = \mathbf{z}'$

otherwise

Set $\mathbf{z}_{n+1} = \mathbf{z}_n$

Note: the chain may contain many duplicated samples due to rejections

- Proof: MH guarantees the detailed balance

Given arbitrary \mathbf{z}_n and \mathbf{z}_{n+1} , if $\mathbf{z}_{n+1} \neq \mathbf{z}_n$, \mathbf{z}_{n+1} must be obtained from accepting a proposal

$$T(\mathbf{z}_n \rightarrow \mathbf{z}_{n+1}) = q(\mathbf{z}_{n+1} | \mathbf{z}_n) \min\left(1, \frac{p(\mathbf{z}_{n+1}, \mathcal{D})q(\mathbf{z}_n | \mathbf{z}_{n+1})}{p(\mathbf{z}_n, \mathcal{D})q(\mathbf{z}_{n+1} | \mathbf{z}_n)}\right)$$

$$= q(\mathbf{z}_{n+1} | \mathbf{z}_n) \min\left(1, \frac{\frac{p(\mathbf{z}_{n+1}, \mathcal{D})/p(\mathcal{D})q(\mathbf{z}_n | \mathbf{z}_{n+1})}{p(\mathbf{z}_n, \mathcal{D})/p(\mathcal{D})q(\mathbf{z}_{n+1} | \mathbf{z}_n)}}\right)$$

$$= q(\mathbf{z}_{n+1} | \mathbf{z}_n) \min\left(1, \frac{p(\mathbf{z}_{n+1} | \mathcal{D})q(\mathbf{z}_n | \mathbf{z}_{n+1})}{p(\mathbf{z}_n | \mathcal{D})q(\mathbf{z}_{n+1} | \mathbf{z}_n)}\right)$$

- Proof: MH guarantees the detailed balance

Given arbitrary \mathbf{z}_n and \mathbf{z}_{n+1} , if $\mathbf{z}_{n+1} \neq \mathbf{z}_n$, \mathbf{z}_{n+1} must be obtained from accepting a proposal

$$T(\mathbf{z}_n \rightarrow \mathbf{z}_{n+1}) = q(\mathbf{z}_{n+1} | \mathbf{z}_n) \min\left(1, \frac{p(\mathbf{z}_{n+1} | \mathcal{D})q(\mathbf{z}_n | \mathbf{z}_{n+1})}{p(\mathbf{z}_n | \mathcal{D})q(\mathbf{z}_{n+1} | \mathbf{z}_n)}\right)$$

$$p(\mathbf{z}_n | \mathcal{D})T(\mathbf{z}_n \rightarrow \mathbf{z}_{n+1}) = p(\mathbf{z}_n | \mathcal{D})q(\mathbf{z}_{n+1} | \mathbf{z}_n) \min\left(1, \frac{p(\mathbf{z}_{n+1} | \mathcal{D})q(\mathbf{z}_n | \mathbf{z}_{n+1})}{p(\mathbf{z}_n | \mathcal{D})q(\mathbf{z}_{n+1} | \mathbf{z}_n)}\right)$$



$$= \min(p(\mathbf{z}_n | \mathcal{D})q(\mathbf{z}_{n+1} | \mathbf{z}_n), p(\mathbf{z}_{n+1} | \mathcal{D})q(\mathbf{z}_n | \mathbf{z}_{n+1}))$$

$$p(\mathbf{z}_{n+1} | \mathcal{D})T(\mathbf{z}_{n+1} \rightarrow \mathbf{z}_n)$$

$$= \min(p(\mathbf{z}_{n+1} | \mathcal{D})q(\mathbf{z}_n | \mathbf{z}_{n+1}), p(\mathbf{z}_n | \mathcal{D})q(\mathbf{z}_{n+1} | \mathbf{z}_n))$$

- Proof: MH guarantees the detailed balance

if $\mathbf{z}_{n+1} = \mathbf{z}_n$

$$T(\mathbf{z}_n \rightarrow \mathbf{z}_{n+1}) = p(\text{reject the proposal}) + \\ p(\text{proposal is } \mathbf{z}_{n+1} \text{ and accept})$$

$$p(\mathbf{z}_n | \mathcal{D}) T(\mathbf{z}_n \rightarrow \mathbf{z}_{n+1}) = p(\mathbf{z}_n | \mathcal{D}) \cdot [p(\text{reject the proposal}) + \\ p(\text{proposal is } \mathbf{z}_{n+1} \text{ and accept})]$$



$$p(\mathbf{z}_{n+1} | \mathcal{D}) T(\mathbf{z}_{n+1} \rightarrow \mathbf{z}_n) = p(\mathbf{z}_n | \mathcal{D}) \cdot [p(\text{reject the proposal}) + \\ p(\text{proposal is } \mathbf{z}_n \text{ and accept})]$$

- If we choose a symmetric proposal distribution

$$q(\mathbf{z}'|\mathbf{z}_n) = q(\mathbf{z}_n|\mathbf{z}') \quad \text{e.g., } \mathcal{N}(\mathbf{z}'|\mathbf{z}_n, \sigma^2 \mathbf{I})$$

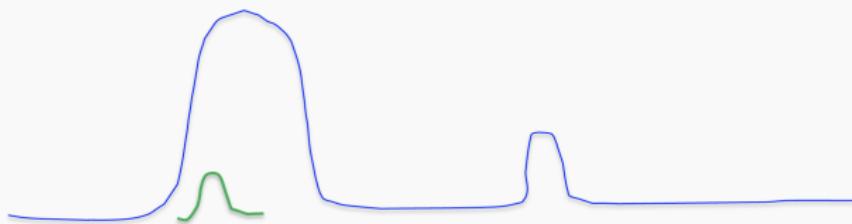
Accept probability: $\min\left(1, \frac{p(\mathbf{z}', \mathcal{D})q(\mathbf{z}_n|\mathbf{z}')} {p(\mathbf{z}_n, \mathcal{D})q(\mathbf{z}'|\mathbf{z}_n)}\right)$

$$= \min\left(1, \frac{p(\mathbf{z}', \mathcal{D})} {p(\mathbf{z}_n, \mathcal{D})}\right)$$

If the proposal increases the model probability, the accept rate is one!

- We need to collect samples that fit the target posterior (e.g., their histogram should be more and more like the posterior). That means, we require many samples on the high-density regions and much less samples on the low-density regions
- However, if the proposals are generated like a random walk through the sample space, a great many proposals will be discarded (due to being in the low-density regions); and much computational cost is wasted

- Take the commonly used Gaussian proposal as an example



- So a key goal to design MCMC algorithms is to reduce random walk behavior!

- General ideas and Markov chain basics
- Metropolis-Hastings algorithm
- Gibbs sampling
- Hybrid Monte-Carlo

- A special type of MH algorithm
- Use conditional distribution to sample each single (or subset of) random variable in the model
- Accept rate is always one
- A good choice when the conditional distribution is tractable and easy to draw samples

$$\mathbf{z} = [z_1, \dots, z_m]^\top \quad p(\mathbf{z}, \mathcal{D}) = p(z_1, \dots, z_m, \mathcal{D})$$

Assume each $p(z_i | \mathbf{z}_{\neg i}, \mathcal{D})$ is tractable and easy to generate samples

$$\mathbf{z}_{\neg i} = [z_1, \dots, z_{i-1}, z_{i+1}, \dots, z_m]^\top$$

- Initialize $\mathbf{z}^{(1)} = [z_1^{(1)}, \dots, z_m^{(1)}]^\top$
- For $t = 1, \dots, T$
 - Sample $z_1^{(n+1)} \sim p(z_1 | z_2^{(n)}, z_3^{(n)}, \dots, z_m^{(n)}, \mathcal{D})$
 - Sample $z_2^{(n+1)} \sim p(z_2 | z_1^{(n+1)}, z_3^{(n)}, \dots, z_m^{(n)}, \mathcal{D})$
 - Sample $z_3^{(n+1)} \sim p(z_3 | z_1^{(n+1)}, z_2^{(n+1)}, \dots, z_m^{(n)}, \mathcal{D})$
 - Sample $z_j^{(n+1)} \sim p(z_j | z_1^{(n+1)}, \dots, z_{j-1}^{(n+1)}, z_{j+1}^{(n)}, \dots, z_m^{(n)}, \mathcal{D})$
 - ...
 - Sample $z_m^{(n+1)} \sim p(z_m | z_1^{(n+1)}, z_2^{(n+1)}, \dots, z_{m-1}^{(n+1)}, \mathcal{D})$

- We can also partition the random variables into sub-vectors, and perform similar alternative sampling

$$\mathbf{z} = [\mathbf{z}_1, \dots, \mathbf{z}_t]^\top$$

$$p(\mathbf{z}_i | \mathbf{z}_1, \dots, \mathbf{z}_{i-1}, \mathbf{z}_{i+1}, \dots, \mathbf{z}_t, \mathcal{D})$$

- This is called block Gibbs sampling

- Matrix factorization

	Movie 1	Movie 2	Movie 3	Movie 4
User 1	3.2	1.2	5	4.0
User 2	2.2	1.0	?	3.0
User 3	2.5	?	4.3	?

	Movie 1	Movie 2	Movie 3	Movie 4
User 1	3.2	1.2	5	4.0
User 2	2.2	1.0	?	3.0
User 3	2.5	?	4.3	?

For each user i , introduce a k -dimensional latent feature vector \mathbf{u}_i

For each movie j , introduce a k -dimensional latent feature vector \mathbf{v}_j

$$p(\mathbf{u}_i) = \mathcal{N}(\mathbf{u}_i | \mathbf{0}, \mathbf{I}) \quad p(\mathbf{v}_j) = \mathcal{N}(\mathbf{v}_j | \mathbf{0}, \mathbf{I})$$

The rating is sampled from a Gaussian

$$p(R_{ij} | \mathbf{U}, \mathbf{V}) = \mathcal{N}(R_{ij} | \mathbf{u}_i^\top \mathbf{v}_j, \tau)$$

	Movie 1	Movie 2	Movie 3	Movie 4
User 1	3.2	1.2	5	4.0
User 2	2.2	1.0	?	3.0
User 3	2.5	?	4.3	?

The joint probability

$$\begin{aligned} & p(\mathbf{U}, \mathbf{V}, \mathbf{R}) \\ &= \prod_i p(\mathbf{u}_i) \prod_j p(\mathbf{v}_j) \prod_{(i,j) \in \mathcal{O}} p(r_{ij} | \mathbf{u}_i^\top \mathbf{v}_j, \tau) \end{aligned}$$

$$\begin{aligned} p(\mathbf{U}, \mathbf{V}, \mathbf{R}) \\ = \prod_i p(\mathbf{u}_i) \prod_j p(\mathbf{v}_j) \prod_{(i,j) \in \mathcal{O}} p(r_{ij} | \mathbf{u}_i^\top \mathbf{v}_j, \tau) \end{aligned}$$

We can use Gibbs sampling to sequentially sample each \mathbf{u}_i and \mathbf{v}_j

The conditional distribution will be Gaussian!

- Proof: the target posterior is invariant to the chain

What is the transition kernel?

$$\begin{aligned} T(\mathbf{z}^{(n)} \rightarrow \mathbf{z}^{(n+1)}) \\ = p(z_1^{(n+1)} | z_2^{(n)}, \dots, z_m^{(n)}, \mathcal{D}) \\ \cdot p(z_2^{(n+1)} | z_1^{(n+1)}, z_3^{(n)}, \dots, z_m^{(n)}, \mathcal{D}) & \quad \text{m steps} \\ \dots \\ \cdot p(z_m^{(n+1)} | z_1^{(n+1)}, z_2^{(n+1)}, \dots, z_{m-1}^{(n+1)}, \mathcal{D}) \end{aligned}$$

- Proof: the target posterior is invariant to the chain

if $\mathbf{z}^{(n)} \sim p(\mathbf{z}|\mathcal{D})$ respect the target posterior

$$T(\mathbf{z}^{(n)} \rightarrow \mathbf{z}^{(n+1)})$$

$$= p(z_1^{(n+1)} | z_2^{(n)}, \dots, z_m^{(n)}, \mathcal{D})$$

$$[z_1^{(n+1)}, z_2^{(n)}, \dots, z_m^{(n)}]^\top$$

$$\cdot p(z_2^{(n+1)} | z_1^{(n+1)}, z_3^{(n)}, \dots, z_m^{(n)}, \mathcal{D})$$

$$[z_1^{(n+1)}, z_2^{(n+1)}, z_3^{(n)}, \dots, z_m^{(n)}]^\top$$

...

...

$$\cdot p(z_m^{(n+1)} | z_1^{(n+1)}, z_2^{(n+1)}, \dots, z_{m-1}^{(n+1)}, \mathcal{D})$$

$$[z_1^{(n+1)}, \dots, z_m^{(n+1)}]$$

$$\mathbf{z}^{(n+1)}$$

- Note that you need also to ensure ergodicity
- A **sufficient** condition is that **none of the conditional distributions be zero anywhere in the sample space** (not hard for continuous distributions)
- If the sufficient condition is **NOT** satisfied, you must explicitly prove the ergodicity!

- One iteration of Gibbs sampling is equivalent to m steps of MH updates, each step with accept prob. 1
- Let us look at one step, w.l.o.g., sample the first element (the other elements are the same)

- Let us look at one step, w.l.o.g., sampling the first element (sampling the other elements are the same)

$$\mathbf{z}_n = [z_1^{(n)}, z_2^{(n)}, \dots, z_m^{(n)}]^\top \longrightarrow \mathbf{z}' = [z_1^{(n+1)}, z_2^{(n)}, \dots, z_m^{(n)}]^\top$$

divide $p(z_2^{(n)}, \dots, z_m^{(n)}, \mathcal{D})$

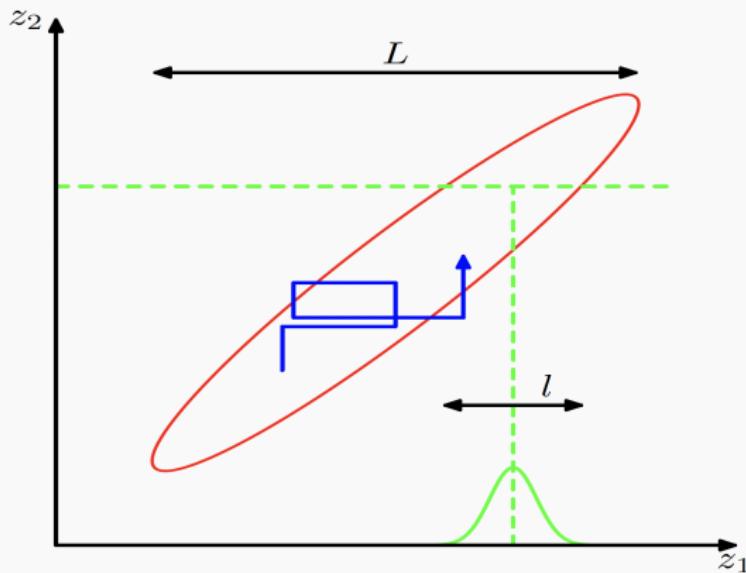
Acceptance probability

$$\min \left(1, \frac{p(z_1^{(n+1)}, z_2^{(n)}, \dots, z_m^{(n)}, \mathcal{D}) p(z_1^{(n)} | z_2^{(n)}, \dots, z_m^{(n)}, \mathcal{D})}{p(z_1^{(n)}, z_2^{(n)}, \dots, z_m^{(n)}, \mathcal{D}) p(z_1^{(n+1)} | z_2^{(n)}, \dots, z_m^{(n)}, \mathcal{D})} \right)$$

||

$$\min \left(1, \frac{p(z_1^{(n+1)} | z_2^{(n)}, \dots, z_m^{(n)}, \mathcal{D}) p(z_1^{(n)} | z_2^{(n)}, \dots, z_m^{(n)}, \mathcal{D})}{p(z_1^{(n)} | z_2^{(n)}, \dots, z_m^{(n)}, \mathcal{D}) p(z_1^{(n+1)} | z_2^{(n)}, \dots, z_m^{(n)}, \mathcal{D})} \right) = 1$$

- Although Gibbs sampling won't reject samples, it may still suffer from inefficient exploration due to strong correlations

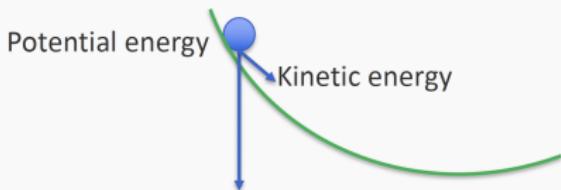


- General ideas and Markov chain basics
- Metropolis-Hastings algorithm
- Gibbs sampling
- Hybrid Monte-Carlo

- Random walk behavior --- waste a lot of samples
- High correlation between different RVs --- slow exploration
- Can we address both problems?

- Also called Hamiltonian MC
- An augmented approach
- Turn the probability to the energy of a physical system
- Augment with other physical properties
- Use the evolution of the physical system (usually described by a set of partial/ordinary differential equations)
- Theoretically can explore the sample space more efficiently, acceptance prob = 1
- Practically limited by the numerical integration error.

- Consider a small ball in a m -dimensional space, without any friction
- Given an initial position and momentum, how does the ball move?



- Characterize how the system evolves
- $\mathbf{z}(t)$: position vector at time t
- Potential energy: $U(\mathbf{z}(t))$
- $\mathbf{r}(t)$: momentum vector at time t
- Kinetic energy: $K(\mathbf{r}(t))$
- Total energy : $H(\mathbf{z}, \mathbf{r}) = U(\mathbf{z}) + K(\mathbf{r})$

- $\mathbf{z}(t)$: position vector at time t
- Potential energy: $U(\mathbf{z}(t))$
- $\mathbf{r}(t)$: momentum vector at time t
- Kinetic energy: $K(\mathbf{r}(t))$
- Total energy : $H(\mathbf{z}, \mathbf{r}) = U(\mathbf{z}) + K(\mathbf{r})$

Evolving:

$$\frac{dz_i}{dt} = \frac{\partial H}{\partial r_i} \quad \mathbf{z} = [z_1, \dots, z_m]^\top$$
$$\frac{dr_i}{dt} = -\frac{\partial H}{\partial z_i} \quad \mathbf{r} = [r_1, \dots, r_m]^\top$$

- How to map our probabilistic model into the system?

$$p(\mathbf{z}, \mathcal{D}) = p(z_1, \dots, z_m, \mathcal{D})$$

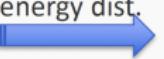
- We take

$$U(\mathbf{z}) = -\log(p(\mathbf{z}, \mathcal{D}))$$

$$K(\mathbf{r}) = \frac{1}{2} \mathbf{r}^\top \mathbf{M}^{-1} \mathbf{r}$$

often takes identity/diagonal matrix

$$H(\mathbf{z}, \mathbf{r}) = U(\mathbf{z}) + K(\mathbf{r})$$

 energy dist.

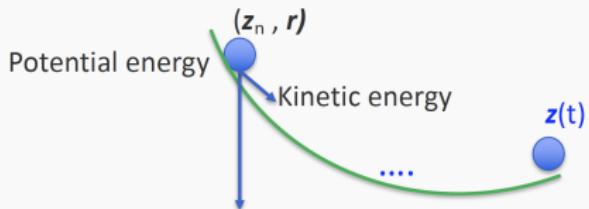
$$p(\mathbf{z}, \mathbf{r}) \propto \exp(-H(\mathbf{z}, \mathbf{r}))$$

What does it include?

$$\left. \begin{array}{l} U(\mathbf{z}) = -\log(p(\mathbf{z}, \mathcal{D})) \\ K(\mathbf{r}) = \frac{1}{2}\mathbf{r}^\top \mathbf{M}^{-1} \mathbf{r} \end{array} \right\} H(\mathbf{z}, \mathbf{r}) = U(\mathbf{z}) + K(\mathbf{r})$$

$$\begin{array}{ll} \frac{dz_i}{dt} = \frac{\partial H}{\partial r_i} & \xrightarrow{\hspace{1cm}} \frac{dz_i}{dt} = [\mathbf{M}^{-1} \mathbf{r}]_i \\ \frac{dr_i}{dt} = -\frac{\partial H}{\partial z_i} & \frac{dr_i}{dt} = -\frac{\partial U}{\partial z_i} \end{array}$$

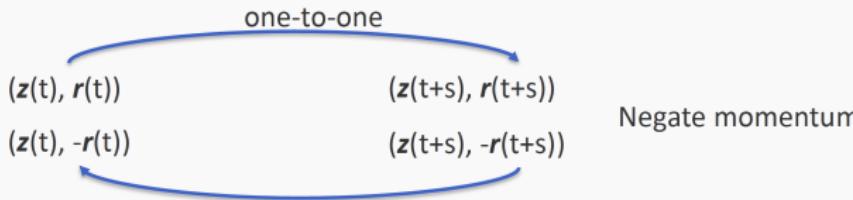
- The key idea: use the current sample z_n and random sample of r , as the **initial state** of the Hamiltonian system; and then **evolve the system to a time t** , pick the $z(t)$ as the proposal and test whether to accept it as z_{n+1}



Note: the proposal is not randomly generated; it is generated deterministically.

- Nice properties to guarantee the detailed balance

1. Reversibility:



Why is it important?

$$p^*(\mathbf{z})T(\mathbf{z} \rightarrow \mathbf{z}') = p^*(\mathbf{z}')T(\mathbf{z}' \rightarrow \mathbf{z})$$

Rigorously speaking, we need to first evolve the system, and then negate the momentum to obtain the new proposal

Now T is a delta function, we need to be able to jump back!

- Nice properties to guarantee the detailed balance

2. Conservation: $\frac{dH}{dt} = 0$ Totally energy does not change

3. Volume preservation: Determinant of Jacobian is always 1



$$\text{For any } t, s \geq 0: \left| \det \frac{\partial [\mathbf{u}(t+s), \mathbf{r}(t+s)]^\top}{\partial [\mathbf{u}(t), \mathbf{r}(t)]^\top} \right| = 1$$

Volume does not change after transformation

Consider an arbitrary dynamic system Ψ_t

Let $\mathbf{v}=(z,r)$ be the extended variable. Define $\mathbf{v}' = \Psi_t(\mathbf{v})$

If the following conditions are satisfied:

- Ψ_t is reversible under R , i.e., $\mathbf{v} = \Psi_t^{-1}(\mathbf{v}') = R(\Psi_t(R(\mathbf{v}')))$
- R is an involution, i.e., $R \circ R(\mathbf{x}) = \mathbf{x}$
- The proposed sample $R(\mathbf{v}')$ is accepted with prob.

$$\min\left\{1, \frac{p(R(\mathbf{v}'))}{p(\mathbf{v})} \left| \det \frac{\partial R \circ \Psi_t(\mathbf{v})}{\partial \mathbf{v}} \right| \right\} \quad \text{otherwise keep } \mathbf{v}$$

Then $p(\mathbf{v})$ is stationary distribution of the Markov chain generated by this Ψ_t and accept test

Consider an arbitrary dynamic system Ψ_t

Let $\mathbf{v} = (\mathbf{z}, \mathbf{r})$ be the extended variable. Define $\mathbf{v}' = \Psi_t(\mathbf{v})$

If the following conditions are satisfied:

- Ψ_t is reversible under R , i.e., $\mathbf{v} = \Psi_t^{-1}(\mathbf{v}') = R(\Psi_t(R(\mathbf{v}')))$
- R is an involution, i.e., $R \circ R(\mathbf{x}) = \mathbf{x}$ R is negating the momentum
- The proposed sample $R(\mathbf{v}')$ is accepted with prob.

conservation $\min\left\{1, \frac{p(R(\mathbf{v}'))}{p(\mathbf{v})} \left| \det \frac{\partial R \circ \Psi_t(\mathbf{v})}{\partial \mathbf{v}} \right| \right\}$ otherwise keep \mathbf{v} volume preservation

Then $p(\mathbf{v})$ is stationary distribution of the Markov chain generated by this Ψ_t and accept test

Energy dist.

Apply the theorem to Hamiltonian system, the accept rate is always 1

$$\left. \begin{array}{l} U(\mathbf{z}) = -\log(p(\mathbf{z}, \mathcal{D})) \\ K(\mathbf{r}) = \frac{1}{2}\mathbf{r}^\top \mathbf{M}^{-1} \mathbf{r} \end{array} \right\} H(\mathbf{z}, \mathbf{r}) = U(\mathbf{z}) + K(\mathbf{r})$$

$$\begin{array}{ll} \frac{dz_i}{dt} = \frac{\partial H}{\partial r_i} & \xrightarrow{\hspace{1cm}} \frac{dz_i}{dt} = [\mathbf{M}^{-1} \mathbf{r}]_i \\ \frac{dr_i}{dt} = -\frac{\partial H}{\partial z_i} & \frac{dr_i}{dt} = -\frac{\partial U}{\partial z_i} \end{array}$$

$$\begin{aligned}\frac{dz_i}{dt} &= [\mathbf{M}^{-1} \mathbf{r}]_i && \text{In practice we often choose} \\ \frac{dr_i}{dt} &= -\frac{\partial U}{\partial z_i} && \mathbf{M} = \text{diag}[s_1, \dots, s_m]\end{aligned}$$

Euler's method: choose step size ϵ , and # of step size L

$$\begin{aligned}r_i(t + \epsilon) &= r_i(t) + \epsilon \frac{dr_i(t)}{dt} = r_i(t) - \epsilon \frac{\partial U(\mathbf{z}(t))}{\partial z_i} && \text{Log joint probability} \\ z_i(t + \epsilon) &= z_i(t) + \epsilon \frac{dz_i(t)}{dt} = z_i(t) + \epsilon \frac{r_i(t)}{s_i}\end{aligned}$$

- Euler's method is a first-order method $O(\epsilon)$
- In practice, people choose Leapfrog method, a second-order method $O(\epsilon^2)$

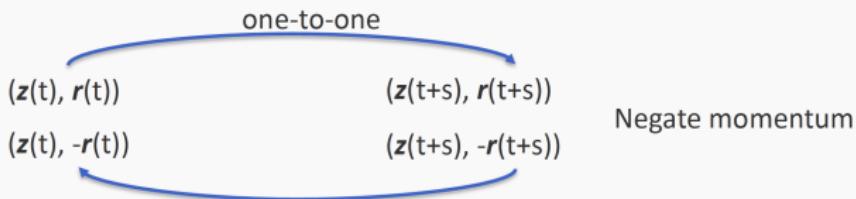
$$r_i(t + \epsilon/2) = r_i(t) - (\epsilon/2) \frac{\partial U(\mathbf{z})}{\partial z_i}$$

$$z_i(t + \epsilon) = z_i(t) + \epsilon \frac{r_i(t + \epsilon/2)}{s_i} \quad \text{introduce half-step}$$

$$r_i(t + \epsilon) = r_i(t + \epsilon/2) - (\epsilon/2) \frac{\partial U(\mathbf{z}(t + \epsilon))}{\partial z_i}$$

- Key properties

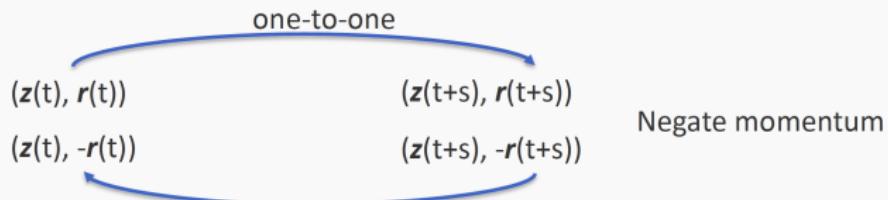
- Reversibility under momentum negation



- Volume preservation: each leap-frog step is a shear transformation and preserves volumes

Question: does conservation still hold?

- Key properties
 - Reversibility under momentum negation



- Volume preservation: each leap-frog step is a shear transformation and preserves volumes

Question: does conservation still hold?

No, because it is a numerical approximation!

Consider an arbitrary dynamic system Ψ_t

Leapfrog

Let $\mathbf{v} = (z, r)$ be the extended variable. Define $\mathbf{v}' = \Psi_t(\mathbf{v})$

If the following conditions are satisfied:

- Ψ_t is reversible under R , i.e., $\mathbf{v} = \Psi_t^{-1}(\mathbf{v}') = R(\Psi_t(R(\mathbf{v}')))$
- R is an involution, i.e., $R \circ R(\mathbf{x}) = \mathbf{x}$ R: momentum negation
- The proposed sample $R(\mathbf{v}')$ is accepted with prob.

$$\min\left\{1, \frac{p(R(\mathbf{v}'))}{p(\mathbf{v})} \left| \det \frac{\partial R \circ \Psi_t(\mathbf{v})}{\partial \mathbf{v}} \right| \right\} \quad \text{otherwise keep } \mathbf{v}$$

Then $p(\mathbf{v})$ is stationary distribution of the Markov chain generated by this Ψ_t and accept test

Note that: due to the numerical error, the accept rate is not guaranteed to be 1

- We augment the latent variable \mathbf{z} , with momentum variables \mathbf{r}
- Construct energy distribution

$$U(\mathbf{z}) = -\log(p(\mathbf{z}, \mathcal{D})) \quad K(\mathbf{r}) = \frac{1}{2}\mathbf{r}^\top \mathbf{M}^{-1} \mathbf{r}$$

$$H(\mathbf{z}, \mathbf{r}) = U(\mathbf{z}) + K(\mathbf{r})$$

$$p(\mathbf{z}, \mathbf{r}) \propto \exp(-H(\mathbf{z}, \mathbf{r}))$$

- We construct a MC to generate samples from $p(\mathbf{z}, \mathbf{r})$

- Step 1: generate new sample for r

$$r_i \sim \mathcal{N}(r_i|0, s_i)$$

(This is a Gibbs sampling step, why? Because the r and z are independent!)

- Step 2: start with current (z, r) and run Leap-frog for L steps with step size ϵ , obtain (z', r') , set $r' = -r'$, accept z' with probability

$$\min\{1, \exp(-H(z', r') + H(z, r))\} = \min\{1, \exp(-U(z') - K(r') + U(z) + K(r))\}$$

otherwise keep z

(This is a Metropolis-hasting step)

- Repeat Step 1 & 2 until get all the samples after burn-in

- Combining multiple Metropolis-hasting steps still yields one valid MH step, so the target posterior is invariant to the transitional kernel of the chain
- Ergodicity: typically satisfied because any value can be sampled from the momentum; only failed when the Leapfrog will produce periodicity; we can overcome this issue by randomly choosing ϵ and L routinely.

- Apply to continuous distributions only, because Leapfrog needs the gradient information
- Very powerful MCMC algorithms.
- Usually much better than original Metropolis Hasting
- Gold-standard for inference in Bayesian neural networks.

- There is a trade-off for the choice (ϵ, L) in the Leapfrog

$$\min\{1, \exp(-H(\mathbf{z}', \mathbf{r}') + H(\mathbf{z}, \mathbf{r}))\}$$

- Large ϵ and L will allow you to explore the space further away, but increase the numerical error and lower the acceptance rate
- Small ϵ and L will be more accurate and so the acceptance rate increases, but the generated samples are not distant.
- In practice, it is very important to tune the two parameters!

- Basic idea of MCMC
- Key concepts: transitional kernel, stationary/invariant/equilibrium distribution, detailed balance...
- Metropolis Hasting and random walk behavior
- Gibbs sampling
- Hybrid Monte-Carlo sampling
- You should be able to implement these algorithms!